

Big Data Analytics at the MPCDF: GPU Crystallography with Python

Giuseppe Di Bernardo

MAX PLANCK COMPUTING AND DATA FACILITY ([MPCDF](#))

giuseppe.di-bernardo@mpcdf.mpg.de

 [@jose_dibernardo](#)

July 7, 2017



Outline

1 X-ray Crystallography: a short overview

2 Big Data Analytics at MPCDF

3 Visualization of Crystal Nano-Structure

Outline

1 X-ray Crystallography: a short overview

- Atom Probe Tomography
- Crystallography: implementation on a GPU

2 Big Data Analytics at MPCDF

- PyNX: A Python-based approach to GPU computing
- PyCUDA

3 Visualization of Crystal Nano-Structure

- Python Plotting for Exploratory Data Analysis
- Getting started with ParaView

Atom Probe Tomography (APT) Microscopy

- Chemical analysis of 3D Volumes
- Atomic scale
- No need of composition calibrations
- Composition of interfaces and defects (dislocations, clusters ...)
- Can be directly compared to 3D simulations at the atomic scale

Figure: Ga grain boundary segregation in a nanocrystalline Al thin-film



Scattering computing from an atomistic model

Theory: X-ray and neutron scattering

- Scattering vector
- Scattering density, at position \mathbf{r}
- Scattered amplitude

$$A(\mathbf{S}) = \int_V \rho(\mathbf{r}) \exp(2i\pi \mathbf{S} \cdot \mathbf{r}) dV = FT[\rho(\mathbf{r})]$$

divide et impera

- FFT is very fast: $\mathcal{O}(N \log N)$
- coordinates in **reciprocal space** ((hkl)) are imposed by the electron density grid

Our approach: *Direct calculation*

- $A(\mathbf{s}) = \sum_{j=0}^{N-1} f_j(s) \exp(2i\pi \mathbf{S} \cdot \mathbf{r}_j)$
- $\mathcal{O}(N \cdot N)$

GPU-accelerated analysis of large X-ray data

- 1 Direct calculation in **3D reciprocal space (hkl)**: **atomic positions (\mathbf{r}_i)** & scattering factors (f_i):

- $$A(\mathbf{s}) = \sum_{j=0}^{N-1} f_j(s) e^{2i\pi \mathbf{s} \cdot \mathbf{r}_j}$$



- computing any point in (hkl) from **any structural model!**

- 2 In crystallography usually we have:

- many atoms, e.g. $N_{\text{atoms}} \mathcal{O}(\geq 1 \times 10^8)$
- many reflections (N_{hkl}) in hkl , e.g. $N_{hkl} \mathcal{O}(128 \times 128 \times 128)$

- 3 It requires **massive parallelism**

- $N_{\text{flop}} \approx 10 \times N_{\text{atoms}} \times N_{hkl} : \mathcal{O}(1 \times 10^{14}) \text{ flop}$
- modern GPU: 10^{13} Flop/s, CPU server: 10^{12} Flop/s
- algorithm well suited for GPU computations, scalable across multiple GPUs $\rightarrow \mathcal{O}(\text{minutes})$

GPU-accelerated analysis of large X-ray data

- 1 Direct calculation in **3D reciprocal space (hkl)**: **atomic positions (\mathbf{r}_i)** & scattering factors (f_i):

- $A(\mathbf{s}) = \sum_{j=0}^{N-1} f_j(s) e^{2i\pi \mathbf{s} \cdot \mathbf{r}_j}$ 

- computing any point in (hkl) from **any structural model!**

- 2 In crystallography usually we have:

- many atoms, e.g. $N_{\text{atoms}} \mathcal{O}(\geq 1 \times 10^8)$ 

- many reflections (N_{hkl}) in hkl , e.g. $N_{hkl} \mathcal{O}(128 \times 128 \times 128)$

- 3 It requires **massive parallelism**

- $N_{\text{flop}} \approx 10 \times N_{\text{atoms}} \times N_{hkl} : \mathcal{O}(1 \times 10^{14}) \text{ flop}$ 

- modern GPU: 10^{13} Flop/s, CPU server: 10^{12} Flop/s 

- algorithm well suited for GPU computations, scalable across multiple GPUs $\rightarrow \mathcal{O}(\text{minutes})$

GPU-accelerated analysis of large X-ray data

- 1 Direct calculation in **3D reciprocal space (hkl)**: **atomic positions (\mathbf{r}_i)** & scattering factors (f_i):

- $A(\mathbf{s}) = \sum_{j=0}^{N-1} f_j(s) e^{2i\pi \mathbf{s} \cdot \mathbf{r}_j}$ 

- computing any point in (hkl) from **any structural model!**

- 2 In crystallography usually we have:

- many atoms, e.g. $N_{\text{atoms}} \mathcal{O}(\geq 1 \times 10^8)$ 

- many reflections (N_{hkl}) in hkl , e.g. $N_{hkl} \mathcal{O}(128 \times 128 \times 128)$

- 3 It requires **massive parallelism**

- $N_{\text{flop}} \approx 10 \times N_{\text{atoms}} \times N_{hkl} : \mathcal{O}(1 \times 10^{14}) \text{ flop}$ 

- modern GPU: 10^{13} Flop/s, CPU server: 10^{12} Flop/s 

- algorithm well suited for GPU computations, scalable across multiple GPUs $\rightarrow \mathcal{O}(\text{minutes})$

Outline

1 X-ray Crystallography: a short overview

- Atom Probe Tomography
- Crystallography: implementation on a GPU

2 Big Data Analytics at MPCDF

- PyNX: A Python-based approach to GPU computing
- PyCUDA

3 Visualization of Crystal Nano-Structure

- Python Plotting for Exploratory Data Analysis
- Getting started with ParaView

Big Data Analytics at the MPCDF

desorption

meth

meth

N_{149.24}

C₁₀H₁₅N

C₁₀
C₁₀

H₁₅
H₁₅

149.24

C₁₀



PyNX: Tools for Nano-structures X-tallography

It is an open-source library created by *Vincent Favre-Nicolin* at ESRF¹with the following main modules:

¹European Synchrotron Radiation Facility; favre@esrf.fr

²New: in PyNX 3.1.0, is called `pynx.scattering`

PyNX: Tools for Nano-structures X-tallography

It is an open-source library created by *Vincent Favre-Nicolin* at ESRF¹with the following main modules:

- 1 pynx.gpu²: X-ray scattering computing using graphical processing units, allowing up to 4×10^{11} reflections/atoms/seconds (e.g., 2x nVidia GeForce GTX 980);

¹European Synchrotron Radiation Facility; favre@esrf.fr

²New: in PyNX 3.1.0, is called pynx.scattering

PyNX: Tools for Nano-structures X-tallography

It is an open-source library created by *Vincent Favre-Nicolin* at ESRF¹with the following main modules:

- 1 pynx.gpu²: X-ray scattering computing using graphical processing units, allowing up to 4×10^{11} reflections/atoms/seconds (e.g., 2x nVidia GeForce GTX 980);
- 2 pynx.ptycho: simulation and analysis of experiments using the ptychography technique;

¹European Synchrotron Radiation Facility; favre@esrf.fr

²New: in PyNX 3.1.0, is called pynx.scattering

PyNX: Tools for Nano-structures X-tallography

It is an open-source library created by *Vincent Favre-Nicolin* at ESRF¹with the following main modules:

- 1 pynx.gpu²: X-ray scattering computing using graphical processing units, allowing up to 4×10^{11} reflections/atoms/seconds (e.g., 2x nVidia GeForce GTX 980);
- 2 pynx.ptycho: simulation and analysis of experiments using the ptychography technique;
- 3 pynx.wavefront: X-ray wavefront propagation in the near, far field, or continuous;

¹European Synchrotron Radiation Facility; favre@esrf.fr

²New: in PyNX 3.1.0, is called pynx.scattering

PyNX: Tools for Nano-structures X-tallography

It is an open-source library created by *Vincent Favre-Nicolin* at ESRF¹with the following main modules:

- 1 `pynx.gpu`²: X-ray scattering computing using graphical processing units, allowing up to 4×10^{11} reflections/atoms/seconds (e.g., 2x nVidia GeForce GTX 980);
- 2 `pynx.ptycho`: simulation and analysis of experiments using the ptychography technique;
- 3 `pynx.wavefront`: X-ray wavefront propagation in the near, far field, or continuous;
- 4 `pynx.cdi`: Coherent Diffraction Imaging reconstruction algorithms using GPU.

¹European Synchrotron Radiation Facility; favre@esrf.fr

²New: in PyNX 3.1.0, is called `pynx.scattering`

PyNX: Hardware requirements

PyNX aims to help computing scattering (X-ray or neutrons) maps for atomic structures, especially if they are *distorted or disordered*.

- High-performance computing with GPUs
 - nVidia's CUDA toolkit³ and the pyCUDA library
 - OpenCL⁴ language, along with pyOpenCL library

GPU-accelerated applications

PyNX provides fast parallel computation of *scattering from large assemblies of atoms ($\gg 10^8$ atoms)* and **1D, 2D or 3D** coordinates ($\gg 10^6$ grid points) in *reciprocal latticespace*.



³<https://developer.nvidia.com/cuda-toolkit>

⁴<https://www.khronos.org/opencl/>

PyNX: Software requirements

PyNX supports Python version **2.7, 3.4** and above (**Anaconda distribution** is recommended).

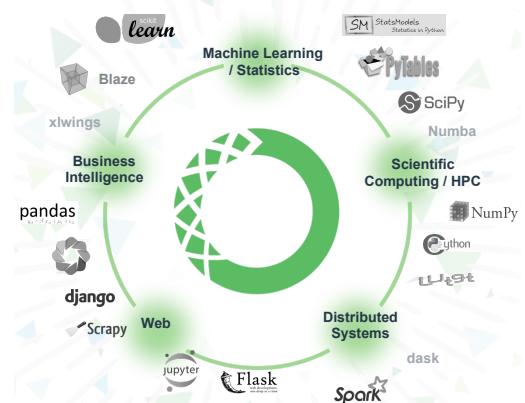
Python interface ⇒ **No need to learn CUDA**

PyNX: Prerequisites

PyNX is available from:

- Mandatory: numpy, matplotlib, pycuda;
- Recommended⁵: pyfftw for cpu calculations;
- Optionally⁶: cctbx library, for *grazing incidence scattering*
\$ conda install -c mx cctbx=20160309

Python is the Common Language



⁶<https://pypi.python.org/pypi/pyFFTW>

⁷<http://cctbx.sourceforge.net/>

PyNX: Software requirements

PyNX supports Python version **2.7, 3.4** and above (**Anaconda distribution** is recommended).

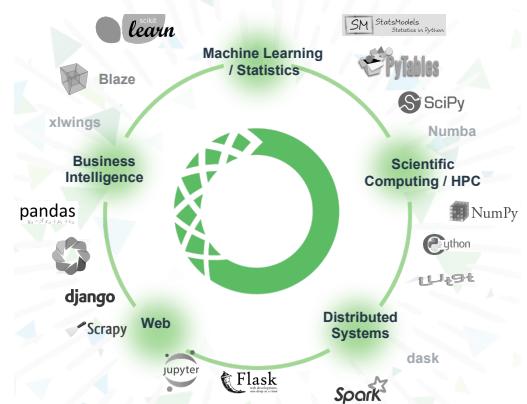
Python interface ⇒ **No need to learn CUDA**

PyNX: Prerequisites

PyNX is available from:

- Mandatory: numpy, matplotlib, pycuda;
- Recommended⁵: pyfftw for cpu calculations;
- Optionally⁶: cctbx library, for *grazing incidence scattering*
\$ conda install -c mx cctbx=20160309

Python is the Common Language



⁶<https://pypi.python.org/pypi/pyFFTW>

⁷<http://cctbx.sourceforge.net/>

PyNX: Software requirements

PyNX supports Python version **2.7, 3.4** and above (**Anaconda distribution** is recommended).

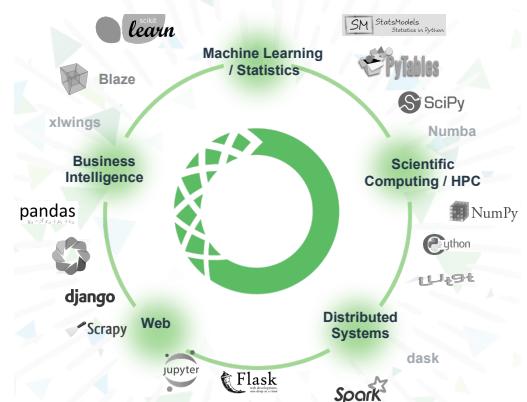
Python interface ⇒ **No need to learn CUDA**

PyNX: Prerequisites

PyNX is available from:

- Mandatory: numpy, matplotlib, pycuda;
- Recommended⁵: pyfftw for cpu calculations;
- Optionally⁶: cctbx library, for *grazing incidence scattering*
\$ conda install -c mx cctbx=20160309

Python is the Common Language



⁶<https://pypi.python.org/pypi/pyFFTW>

⁷<http://cctbx.sourceforge.net/>

PyNX: Software requirements

PyNX supports Python version **2.7, 3.4** and above (**Anaconda distribution** is recommended) .

Python interface ⇒ **No need to learn CUDA**

Install PyNX

PyNX is available from:

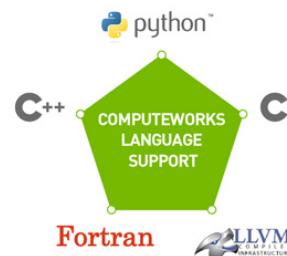
- <http://ftp.esrf.fr/pub/scisoft/PyNX/>

- <http://gitlab.esrf.fr/favre/PyNX>

- PyPI:

```
$ pip install pynx
```

Python is also a great glue language!



PyNX: Software requirements

PyNX supports Python version **2.7, 3.4** and above (**Anaconda distribution** is recommended) .

Python interface ⇒ **No need to learn CUDA**

Install PyNX

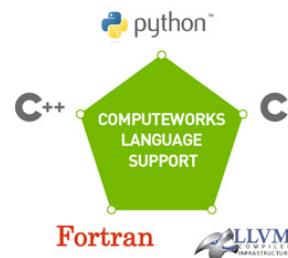
PyNX is available from:

- <http://ftp.esrf.fr/pub/scisoft/PyNX/>
- <http://gitlab.esrf.fr/favre/PyNX>

■ PyPI:

```
$ pip install pynx
```

Python is also a great glue language!



PyNX: Software requirements

PyNX supports Python version **2.7, 3.4** and above (**Anaconda distribution** is recommended) .

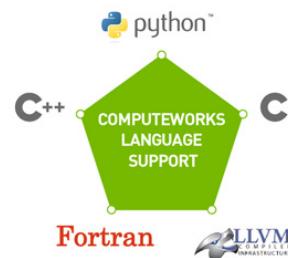
Python interface ⇒ **No need to learn CUDA**

Install PyNX

PyNX is available from:

- <http://ftp.esrf.fr/pub/scisoft/PyNX/>
- <http://gitlab.esrf.fr/favre/PyNX>
- PyPI:
`$ pip install pynx`

Python is also a great glue language!



Why do Scripting for GPUs?

- GPUs are everything that *high-level scripting languages* **are not**.
 - Highly parallel
 - Very architecture-sensitive
 - Built for maximum FP/memory throughput
→ complement each other
- Scripting + GPU : A good combination
 - ↓
- Python + CUDA⁷ = **PyCUDA**
- Python + OpenCL = **PyOpenCL**



Figure: courtesy by
A. Klöckner, Nvidia
GTC 2010

⁸CUDA is a general purpose parallel computing architecture that leverages the parallel compute engine in NVIDIA graphics processing units (GPUs)

But before PyCUDA: why Python at all?

- general purpose 😊
- simple to learn and use 😊
- extensible and embeddable: Python C API
- science oriented too (NumPy, SciPy, mpi4py)
- great visualization tools (Bokeh, Matplotlib)
- very well documented



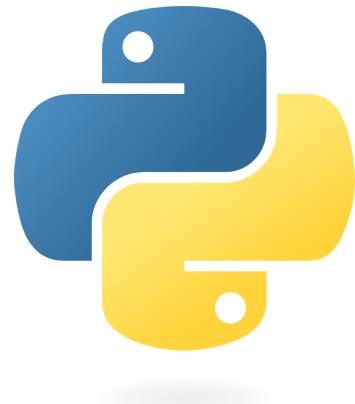
NumPy: package for large, N-dimensional array object 😊

- Vectors, Matrices, ... array computing is fast
- A+B, sin(A), dot(A,B)
- la.solve(A,b), la.eig(A)
- cube[;, :, n-k:n+k], cube+5
- FFT's, tools for integrating C/C++ and Fortran code

Scripting: Python⁸

One example of a scripting language: Python

- Mature
- Large and active community
- Emphasizes readability
- Written in widely-portable C
- A 'multi-paradigm' language
- Rich ecosystem of sci-comp related software



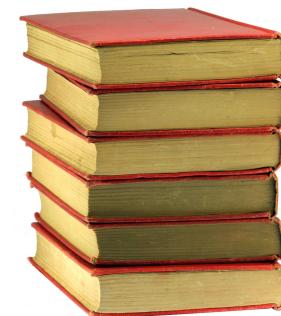
⁸Slide courtesy by A. Klöckner, Nvidia GTC 2010

PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python.

PyCUDA Philosophy: maps all CUDA into Python

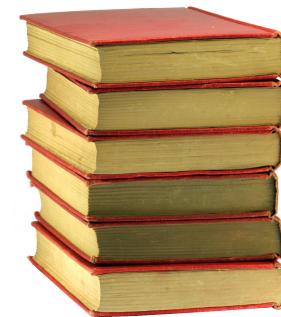
PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**



PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**

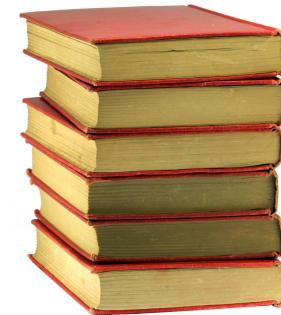
- **Robustness:** automatic management of object lifetimes and error checking



PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**

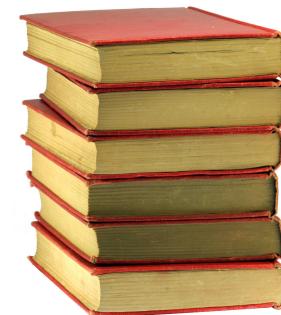
- **Robustness:** automatic management of object lifetimes and error checking
- **Completeness:** automatic error checking



PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**

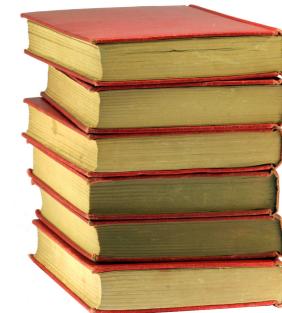
- **Robustness:** automatic management of object lifetimes and error checking
- **Completeness:** automatic error checking
- **Convenience:** provide abstractions (comes with ready-made on-GPU linear algebra, reduction, scan. Add-on packages for FFT and LAPACK available.)



PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**

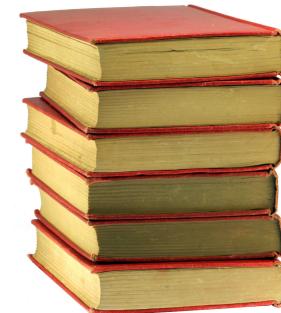
- **Robustness:** automatic management of object lifetimes and error checking
- **Completeness:** automatic error checking
- **Convenience:** provide abstractions (comes with ready-made on-GPU linear algebra, reduction, scan. Add-on packages for FFT and LAPACK available.)
- Integrate tightly with **NumPy**



PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**

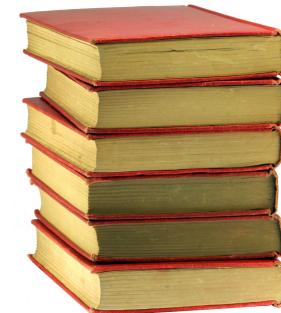
- **Robustness:** automatic management of object lifetimes and error checking
- **Completeness:** automatic error checking
- **Convenience:** provide abstractions (comes with ready-made on-GPU linear algebra, reduction, scan. Add-on packages for FFT and LAPACK available.)
- Integrate tightly with **NumPy**
- **Speed:** PyCUDA's base layer is written in C++ (near-zero wrapping overhead)



PyCUDA Philosophy: maps all CUDA into Python

PyCUDA lets you completely access to Nvidia's CUDA parallel computation API from Python. **Key Features:**

- **Robustness:** automatic management of object lifetimes and error checking
- **Completeness:** automatic error checking
- **Convenience:** provide abstractions (comes with ready-made on-GPU linear algebra, reduction, scan. Add-on packages for FFT and LAPACK available.)
- Integrate tightly with **NumPy**
- **Speed:** PyCUDA's base layer is written in C++ (near-zero wrapping overhead)
- Complete, helpful [documentation](#)



Whetting your appetite: SourceModule

```
import pycuda.driver as cuda
import pycuda.autoinit
import numpy as np
a = np.random.randn(4,4).astype(np.float32)
a_gpu = cuda.mem_alloc(a.nbytes)
# host-to-device
cuda.memcpy_htod(a_gpu, a)

mod = cuda.SourceModule("""
    __global__ void multiply_by_two(float *a)
{
    int idx = threadIdx.x + threadIdx.y*4;
    a[idx] *= 2;
}
""")

func = mod.get_function("multiply_by_two")
func(a_gpu, bloc=(4,4,1))
```

COMPUTE KERNEL

Whetting your appetite: GPUArrays

```
import numpy as np
import pycuda.autoinit
import pycuda.gpuarray as gpuarray

a_gpu = gpuarray.to_gpu(np.random.randn(4,4).astype(np.float32))
a_doubled = (2*a_gpu).get()
print a_doubled
print a_gpu
```

GPUArrays: computational linear algebra

- element-wise algebraic operations (+, -, *, /, sin, cos, exp)
- tight integration with numpy
 - gpuarray.to_gpu(numpy_array)
 - numpy_array = gpuarray.get()
- mixed data types (int32 + float32 = float64)

PLANCK



PyCUDA Workflow: "Edit-Run-Repeat"

A two-fold aim:

- 1 usage of *existing* CUDA C
- 2 *on top of the first layer*, PyCUDA \Rightarrow *abstractions*

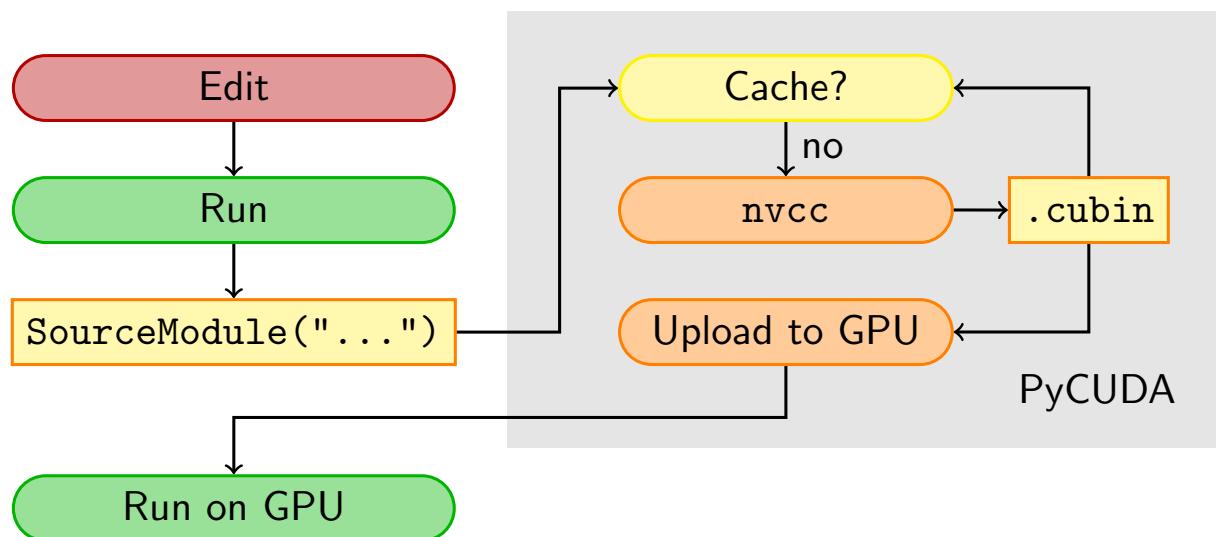


Figure: A. Klöckner et al. 2013, <https://arxiv.org/abs/1304.5553>

PyCUDA: Vital Information

- Availability: Freely down-loadable from this location
 - Open source
 - MIT Licensed
- requires **NumPy**, Python 2.4+ (Win/OS X/Linux)
- Support via mailing list
- For further information see:
 - Main PyCUDA page
 - PyCUDA Wiki
 - PyCUDA FAQ



Figure: courtesy by
A. Klöckner, Nvidia
GTC 2010

PyNX kernel: __global__ void CUDA_fhkl()

```
1 const unsigned long ix=threadIdx.x+blockDim.x*blockIdx.x; // Thread idx
2 const float h=twopi*vh[ix];
3 const float k=twopi*vk[ix];
4 const float l=twopi*vl[ix];
5 float fr=0,fi=0;
6 __shared__ float x[BLOCKSIZE]; // Shared memory between BLOCKSIZE %d
7 __shared__ float y[BLOCKSIZE]; // parallel threads
8 __shared__ float z[BLOCKSIZE];
9 long at=0;
10 for (at<=(natoms-BLOCKSIZE);at+=BLOCKSIZE) {
11     __syncthreads();
12     x[threadIdx.x]=vx[at+threadIdx.x]; // "Coalesced" transfer
13     y[threadIdx.x]=vy[at+threadIdx.x]; // to shared memory
14     z[threadIdx.x]=vz[at+threadIdx.x];
15     __syncthreads();
16     for(unsigned int i=0;i<BLOCKSIZE;i++) { // Each thread computes
17         float s,c; // a single reflection
18         __sincosf(h*x[i] + k*y[i] + l*z[i] , &s,&c);
19         fr +=c; // fast, intrinsic trigonometric function
20         fi +=s;
21     }__syncthreads();}
```

PyNX kernel: __global__ void CUDA_fhkl()

```
1  __syncthreads()
2  ;
3
4  /* Take care of remaining atoms */
5  if(threadIdx.x<(natoms-at)) {
6      x[threadIdx.x]=vx[at+threadIdx.x];
7      y[threadIdx.x]=vy[at+threadIdx.x];
8      z[threadIdx.x]=vz[at+threadIdx.x];
9  }
10 __syncthreads();
11 for(long i=0;i<(natoms-at);i++) {
12     float s,c;
13     __sincosf(h*x[i] + k*y[i] + l*z[i] , &s,&c);
14     fr +=c;
15     fi +=s;
16 }
17 fhkl_real[ix]+=fr;
18 fhkl_imag[ix]+=fi;
19 }
```

PyNX Python interface: loading the data

```
1 import numpy as np
2 import time
3
4 @timeit
5 def do_work(*args, **kwargs):
6     def read_pos(file_name):
7         """Loads an APT.pos file
8             Columns: x, y, z, Da
9         """
10        print("reading in the data")
11        f = open(file_name, mode='rb') # Load orthonormal coord.
12        dt_type = np.dtype({'names':['x', 'y', 'z', 'Da'],
13                            'formats':[ '>f4', '>f4', '>f4', '>f4']})
14        pos = np.fromfile(f, dt_type, -1)
15        f.close()
16        print("the data contain: {0:.5e} atoms".format(pos.size))
17        return pos
```

PyNX Python interface: parsing the data

```
16     if args[1] is not None:  
17         fname = args[1]  
18  
19     df = readpos(fname)  
20     xEl = df['x'] # slicing columns by labels  
21     yEl = df['y']  
22     zEl = df['z']  
23     lattice_parameter = 0.4045  
24     xEl /= lattice_parameter # Convert to fractional coordinates  
25     yEl /= lattice_parameter  
26     zEl /= lattice_parameter  
27     N = 128  
28     h = np.linspace(-1.1, 1.1, num=N, endpoint=True) # HKL as 3D  
29     k = np.linspace(-1.1, 1.1, num=N, endpoint=True)[:, np.  
           newaxis]  
30     l = np.linspace(-1.1, 1.1, num=N, endpoint=True)[:, np.  
           newaxis, np.newaxis]
```

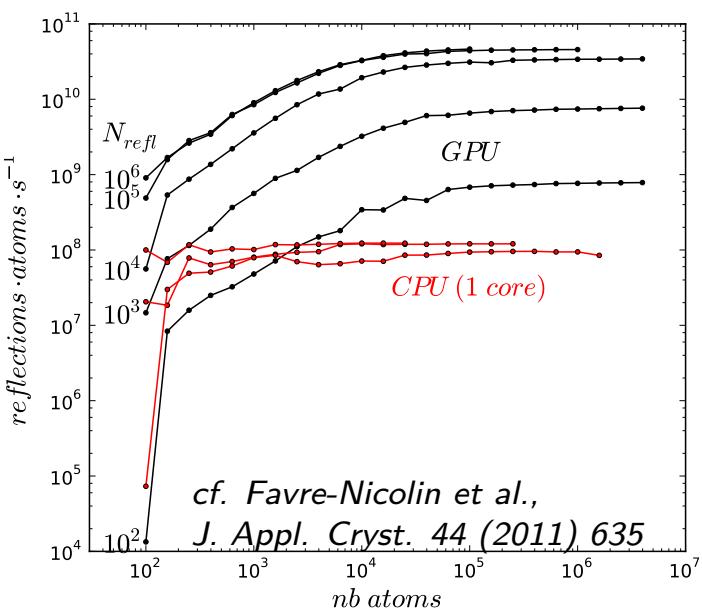
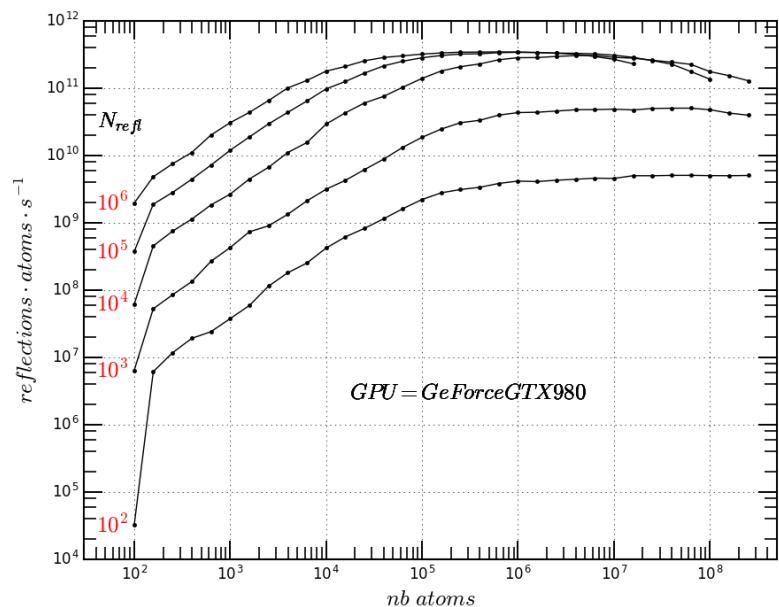
PyNX Python interface: The Fourier Transform

```
31     from pynx import gpu # main module for scattering computation
32     gpu_name = args[0] # identify GPU, e.g. GeForce GTX980
33
34     if args[0] is not None:
35         fhkl, dt = gpu.Fhkl_thread(h, k, l, xEl, yEl, zEl,
36             gpu_name=args[0], verbose=True, language="CUDA",
37             cl_platform="nvidia")
38         print("The FT computed in dt = {0:7.5f}: ".format(dt))
39         print("{0:d} fourier number points".format(fhkl.size))
40
41     return fhkl # a numpy complex64 array
```

gpu.Fhkl_thread

- computes $F(hkl) = \sum_i \exp[2j\pi(x_i \times h + y_i \times k + z_i \times l)]$;
- distributes the scattering maps on **several GPU**.

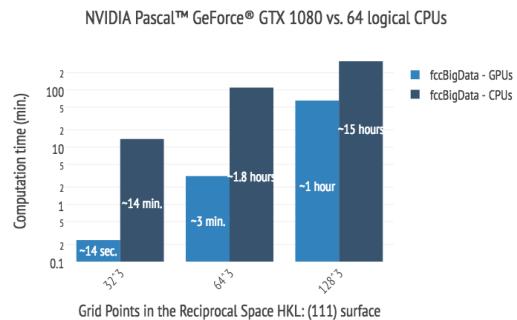
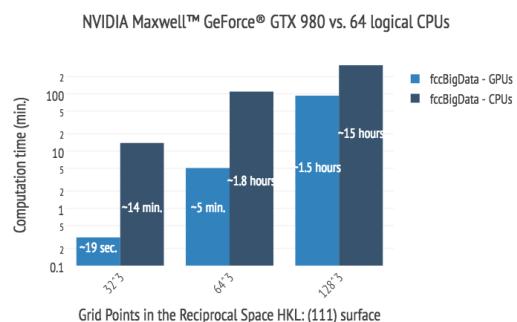
PyNX: Performance



assessment of **GPU-accelerated python** software package PyNX

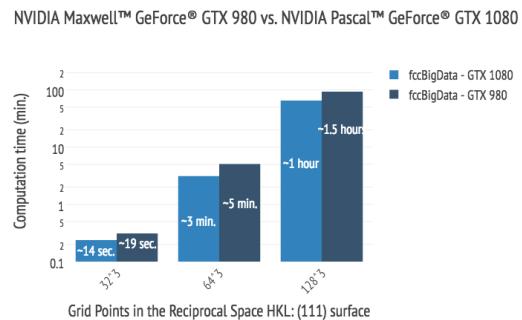
- Effective (single precision) throughput per GPU:
up to 3.5×10^{11} **reflections · atoms · s⁻¹**

PyNX: benchmarks at the MPCDF



PyNX computing time at the MPCDF for 10^8 atoms

- 1 $2 \times$ nVidia GTX 980
- 2 $2 \times$ nVidia GTX 1080
- 3 interactive plots with Plotly



Deploying Data Science Projects - Clusters

```
#!/bin/bash -l
#SBATCH -J pynx_run
# Queue:
#SBATCH -partition=gpu
# Node feature:
#SBATCH -constraint="gpu"
# N. of nodes and MPI tasks per node:
#SBATCH -nodes=1
#SBATCH -ntasks-per-node=32
# wall clock limit:
#SBATCH -time=02:00:00
echo "starting the job ..."
python crystAl.py -g="GeForce GTX 980"
echo "...done"
```

SLURM workload manager: basic usage

- qsub: submit your job;
- qstat: query queue/job status;
- qdel: delete your job

SUPERCOMPUTING DRIVES
SCIENCE THROUGH SIMULATIONS



*The MPG SuperComputer Hydra
at the MPCDF*

Outline

1 X-ray Crystallography: a short overview

- Atom Probe Tomography
- Crystallography: implementation on a GPU

2 Big Data Analytics at MPCDF

- PyNX: A Python-based approach to GPU computing
- PyCUDA

3 Visualization of Crystal Nano-Structure

- Python Plotting for Exploratory Data Analysis
- Getting started with ParaView

Image Processing in Python: Toolbox for SciPy

```

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d.art3d import
Poly3DCollection
from mpl_toolkits.mplot3d import Axes3D
from skimage import measure
pow_spec = abs(np.log10(abs(fhkl)**2))
mean_iso = (pow_spec.max() + pow_spec.min()) / 2.
verts, faces =
measure.marching_cubes(pow_spec, mean_iso)
fig = plt.figure()
ax = plt.axes(projection="3d")
mesh = Poly3DCollection(verts[faces])
ax.add_collection3d(mesh)
plt.show()

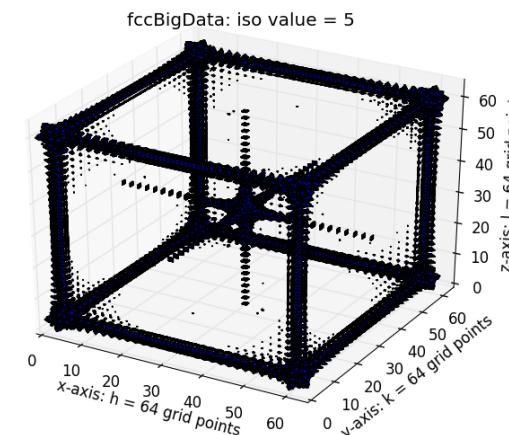
```

Marching Cubes

- 1 MC is an algorithm to extract a **2D surface mesh from a 3D volume**;
- 2 ‘verts[faces]’ to generate a collection of triangles;
- 3 Display triangular mesh using Matplotlib.



scikit-image
image processing in python



NumPy to VTK

```
1 from pyevtk.hl import gridToVTK
2 h=k=l = np.linspace(-1.1, 1.1, num=N, endpoint=True, dtype='
    float64')
3 gridToVTK("/viz/vtk/" + filename, h, k, l, pointData = {"pow_spec
    " : pow_spec})
```

PyEVTK: a great little package by a Paulo Herrera

- 1 \$ pip install pyevtk
- 2 save **NumPy** arrays straight to different types of **VTK XML-based**
- 3 visualize and process your **NumPy** arrays with any of the flagship VTK applications such as **ParaView**, **VisIt**, **Mayavi**

What is ParaView?

- ParaView is an **open-source application** and **architecture** for **display** and analysis of scientific data sets;
- ParaView has been demonstrated to process billions of unstructured cells and to process over a trillion structured cells;
- ParaView's parallel framework has run on over 100,000 processing cores (from notebooks to world's largest supercomputers);
- ParaView's **key features** are:
 - An open-source scalable, multi-platform for visualizing **2D/3D data** (excels at traditional scientific vis qualitative 3D rendering);
 - An extensible, modular **architecture** based on open standards e.g. Custom apps, plugins, Python scripting, ParaViewWeb, Catalyst
 - Support for distributed computation models to process **large data sets**;
 - An open, extensible, and **intuitive user interface**.

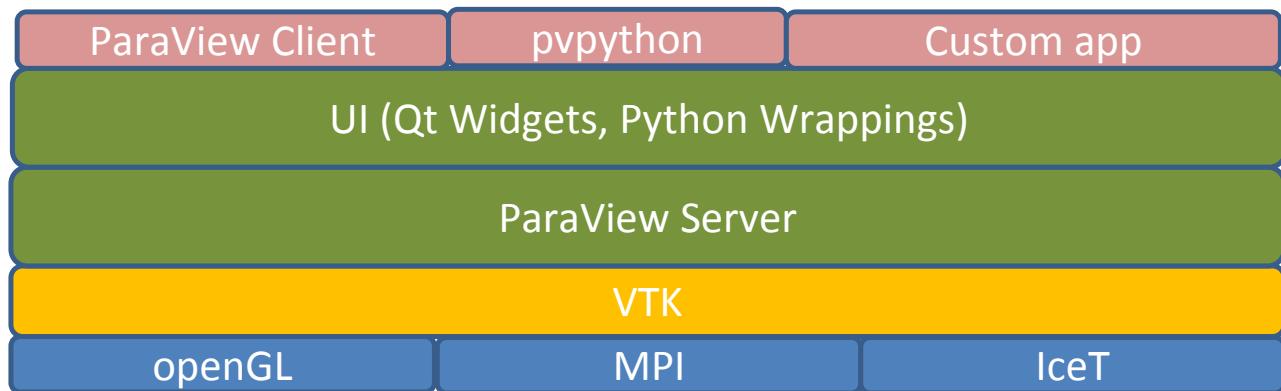
ParaView: a standard de-facto

- ParaView is primarily developed and published by Kitware Inc.
 - A flexible BSD 3-clause license;
 - Commercial maintenance and support;
- ParaView is used by many academic, government, and commercial institutions all over the world;
- ParaView is downloaded roughly 100,000 times every year;
- ParaView also won the HPCwire Readers' Choice Award and HPCwire Editors' Choice Award for Best HPC Visualization Product or Technology.



ParaView: The architecture

The application most people associate with ParaView is really just a small client application built on top of a **tall stack** of libraries that provide ParaView with its functionality.



ParaView: The big picture

The application most people associate with ParaView is really just a small client application built on top of a **tall stack** of libraries that provide ParaView with its functionality.

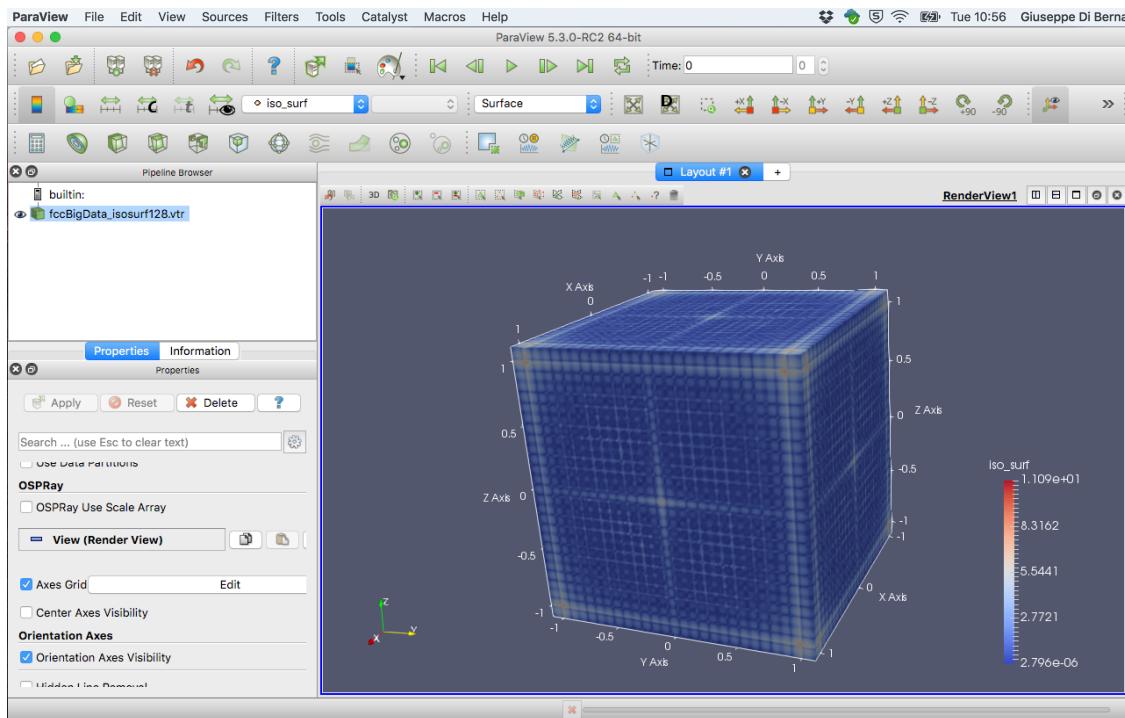
Full Stack

- ParaView comes with a `pvpython` application that allows you to automate the visualization and post-processing with Python scripting;
- ParaView Server library `pvserver` provides the abstraction layer necessary for running parallel, interactive visualization. It relieves the client application from most of the issues concerning if and how ParaView is running in parallel;
- The **Visualization Toolkit** (VTK) provides the basic visualization and rendering algorithms ⇒ **Data-Flow Paradigm**

ParaView: Loading, filtering and rendering

GUI elements

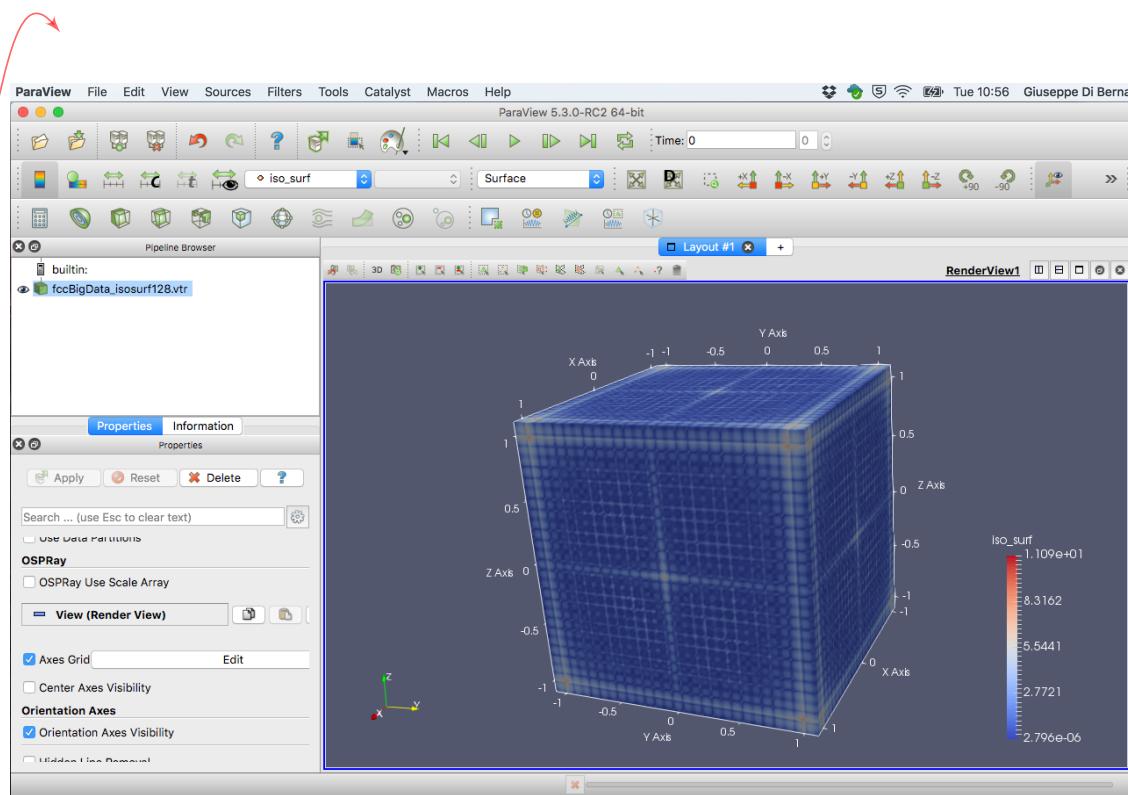
- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



ParaView: Loading, filtering and rendering

GUI elements

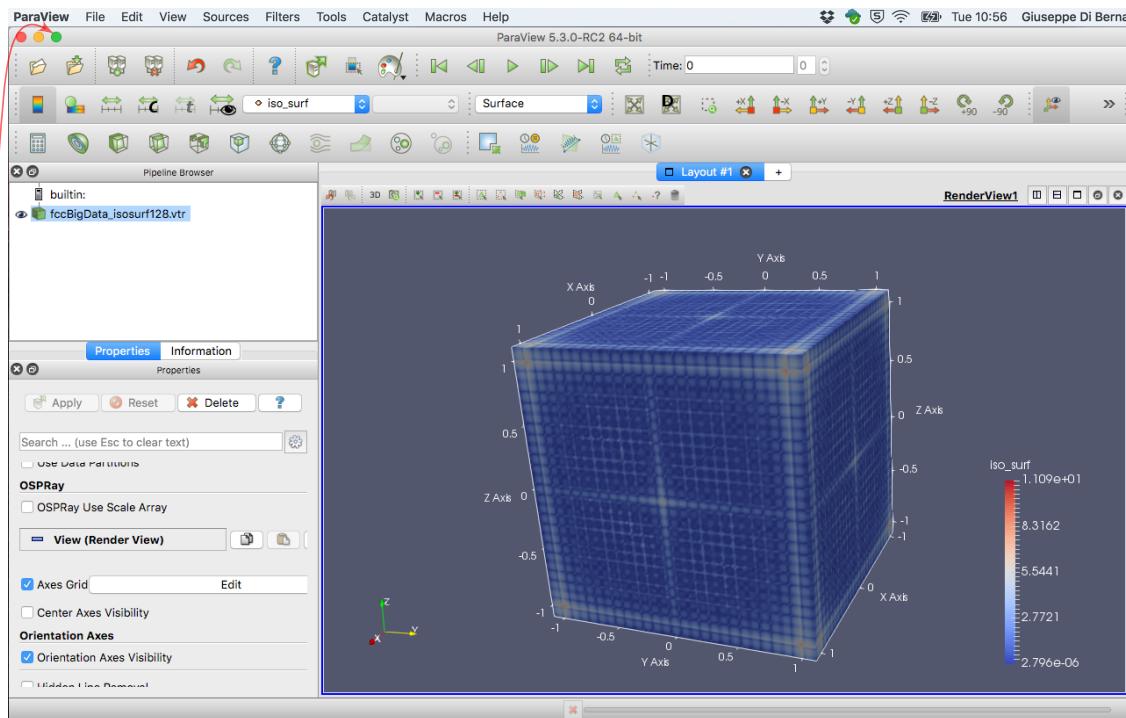
- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



ParaView: Loading, filtering and rendering

GUI elements

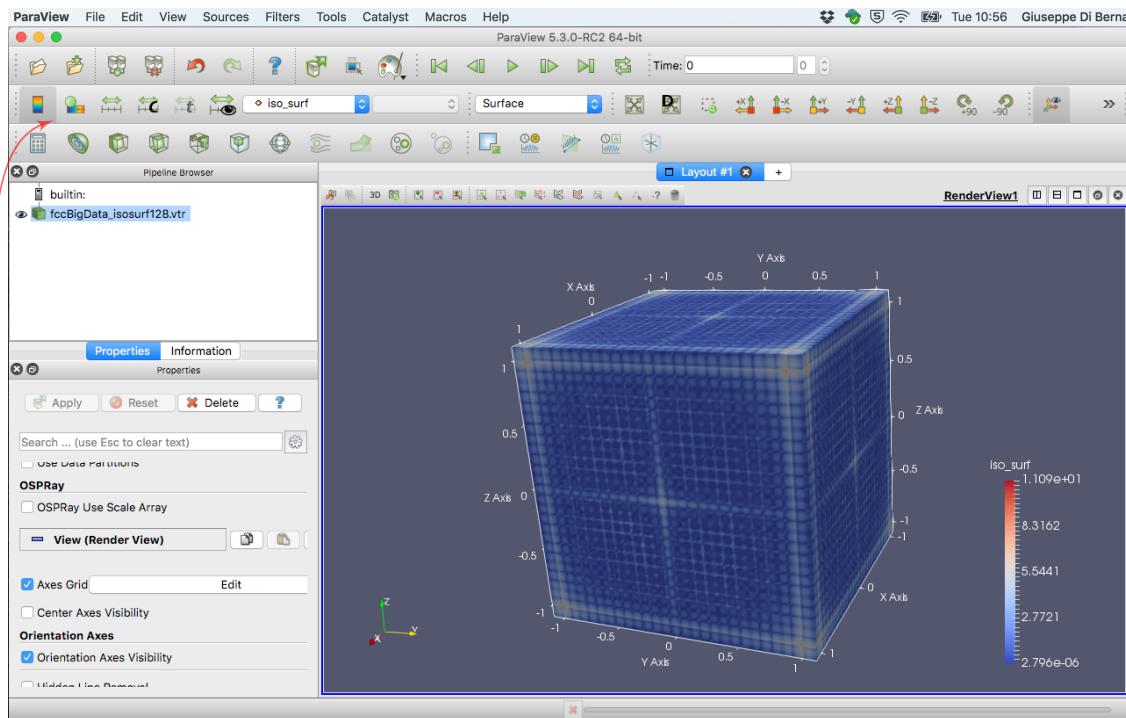
- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



ParaView: Loading, filtering and rendering

GUI elements

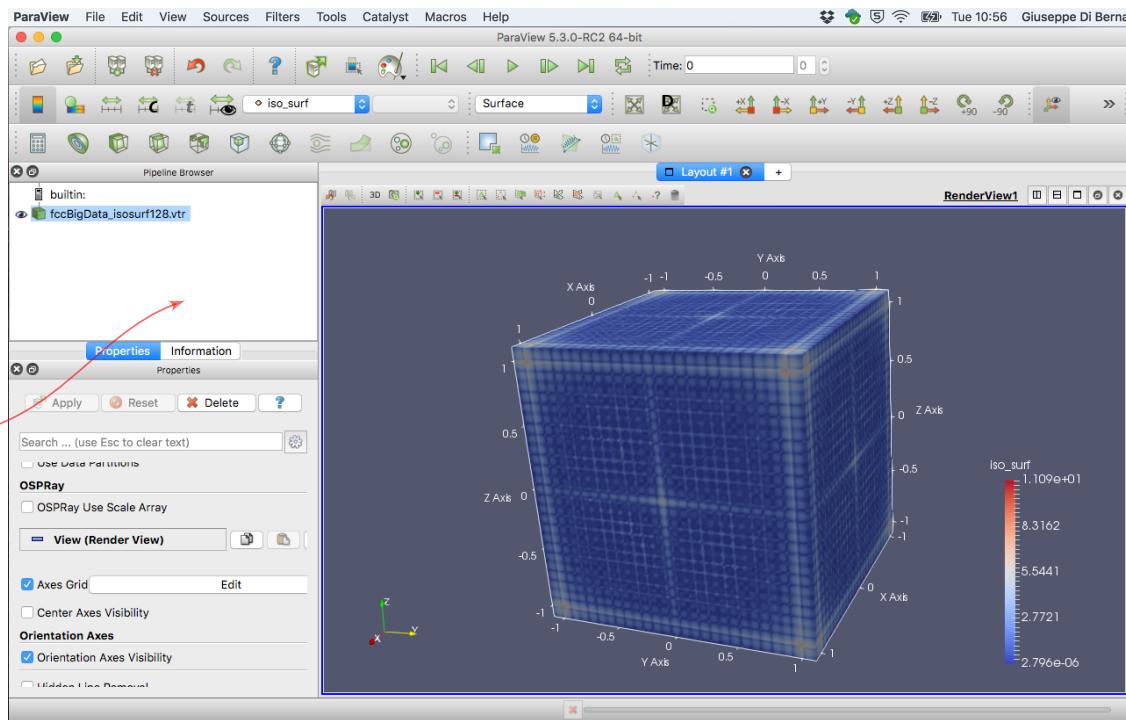
- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



ParaView: Loading, filtering and rendering

GUI elements

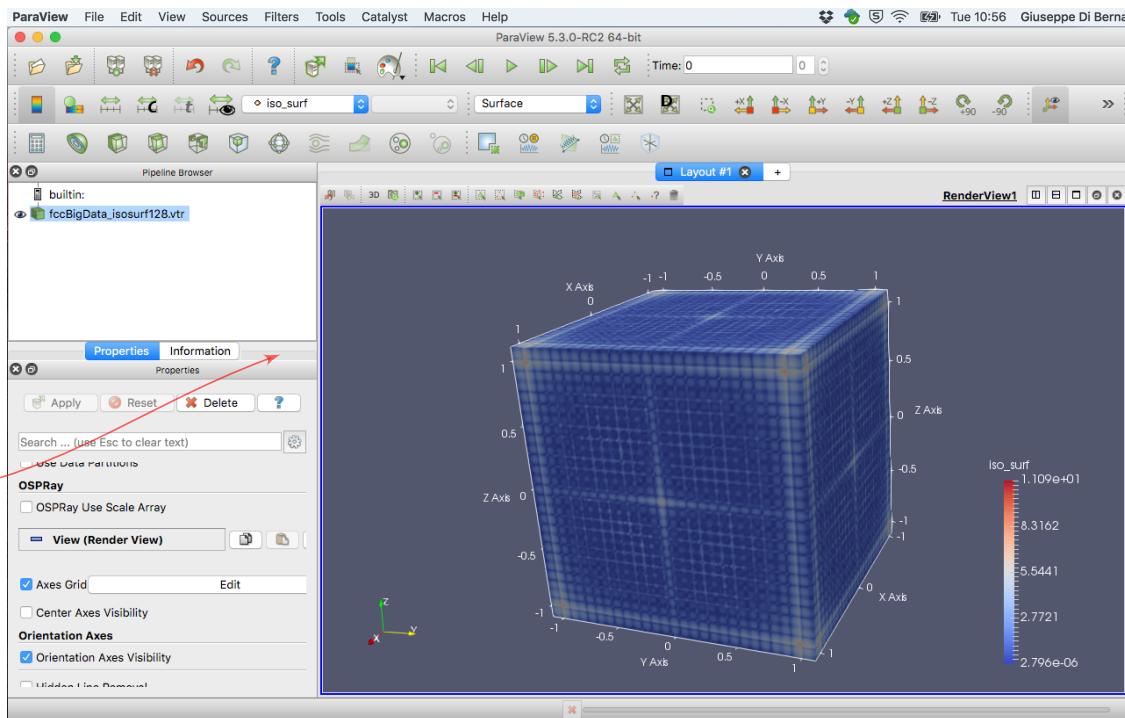
- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



ParaView: Loading, filtering and rendering

GUI elements

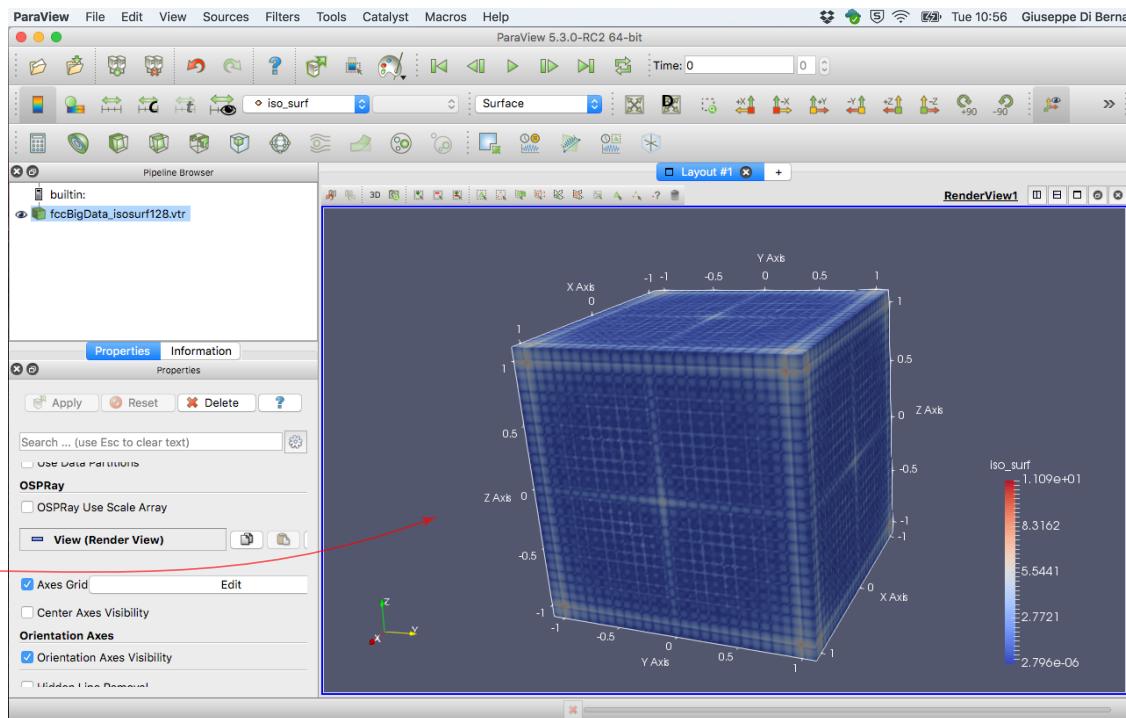
- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



ParaView: Loading, filtering and rendering

GUI elements

- Menu
- Toolbars
- Pipeline
- Inspector
- Help
- Views



GUI elements definition

- 1 Menu Bar:** allows you to access the majority of features;
- 2 Toolbars:** provide quick access to the most commonly used features within ParaView;
- 3 Pipeline Browser:** ParaView manages the reading and iterating of data with a pipeline. A convenient list of pipeline objects with an indentation style that shows the pipeline structure;
- 4 Properties Panel:** to view and change the parameters of the current pipeline object. The properties are by default coupled with an information tab that shows a basic summary of the data produced by the pipeline object;
- 5 3D View:** to present the data so that you may view, interact with, and explore your data.

Filters Menu

...some of them

- 1 **Calculator:** evaluates a user-defined expression on a *per-point* or *per-cell* basis
- 2 **Contour:** extracts the points, curves, or surfaces where a scalar field is equal to a user-defined value. This surface is often also called *iso-surface*
- 3 **Threshold:** extracts cells that lie within a specified range of a scalar-field



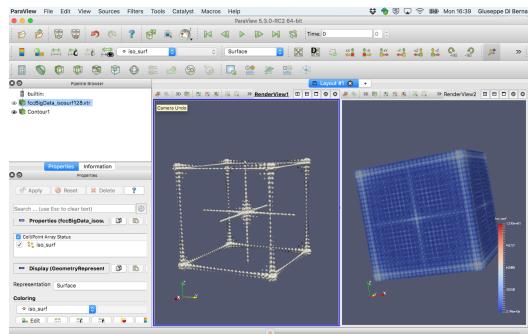
Other Filters

...some of them

- Recent
- AMR (adaptive mesh refinement)
- CTH
- Common (the same list available in the filters toolbars)
- Cosmology

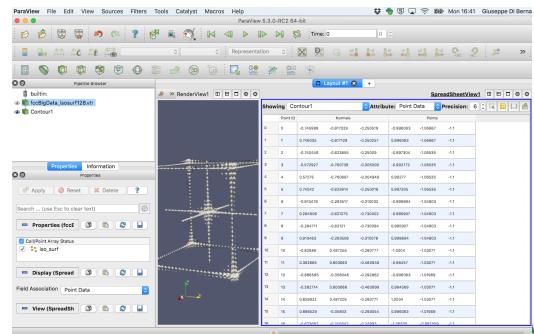
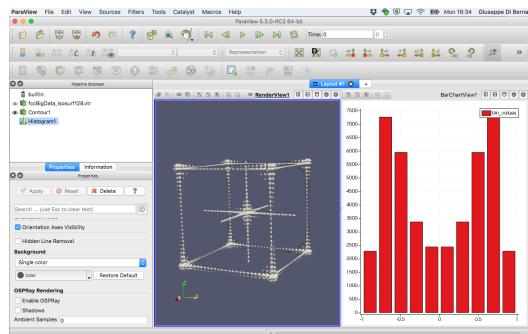
- **Data Analysis:** quantitative values from the data
- Material Analysis
- **Statistics:** descriptive statistics of the data (primarily in tabular form)
- Temporal
- Alphabetical

ParaView: Multiple views



Some examples

- 1 3d + 3d
- 2 3d + Histograms
- 3 3d + table

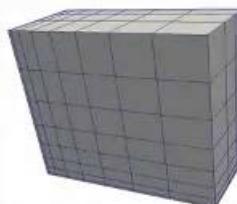


ParaView: The dataset types (sampling structures)

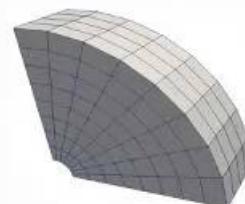
ParaView was designed primarily to handle data with spatial representation. Thus the **data types** used in ParaView are meshes



Uniform Rectilinear
(vtkImageData)



Non-Uniform Rectilinear
(vtkRectilinearData)



Curvilinear
(vtkStructuredData)



Polyhedral
(vtkPolyData)



Unstructured Grid
(vtkUnstructuredGrid)

Multi-block

Hierarchical Adaptive
Mesh Refinement
(AMR)

Hierarchical Uniform
AMR
Octree

Data Selection

The goal of visualization is often to find the important details within a large body of information. ParaView's selection abstraction is an important simplification of this process.

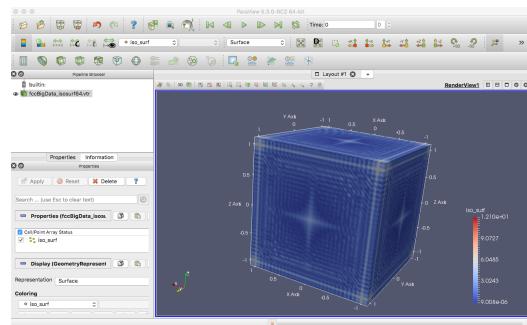
Selection is the act of identifying a subset of some dataset. More specifically the subset identifies particular select points, cells, or blocks within any single data set.

In ParaView, selection can take place at any time, and the program maintains a current selected set that is linked between all views.

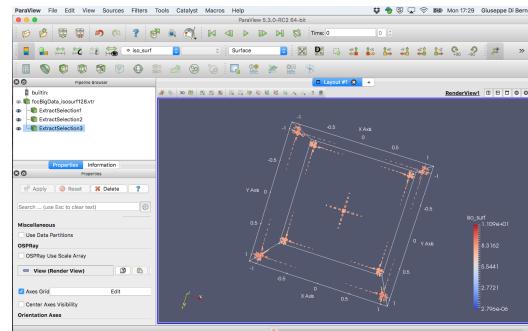
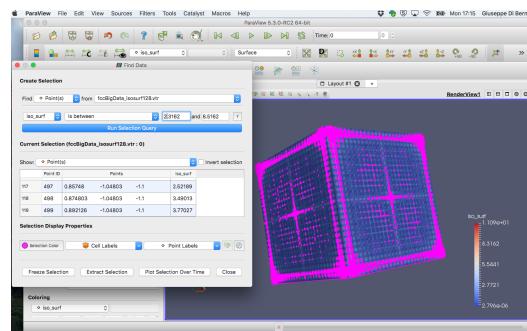
- 1 The most direct means to create a selection is via the Find Data dialog 
- 2 Another way of creating a selection is to pick elements right inside the 3D view: **rubber-band selection**

Query Data by Attribute Values - Find Data dialog

Some examples



- 1 load, filter and display your volume data
- 2 find data matching your criteria
- 3 extract selection (Run Selection Query)



Query Data Spatially - Selection

...by clicking and dragging the mouse in the 3D view

- **Select Cells On (Surface):**

cells that are visible in the view under a rubber band

- **Select Points On (Surface):** points that are visible in the view under a rubber band

- **Select Cells Through (Frustum):** cells that exist under a rubber band

- **Select Points Through (Frustum)**

- **Select Cells with Polygon:**

you draw a polygon

- **Select Points with Polygon**

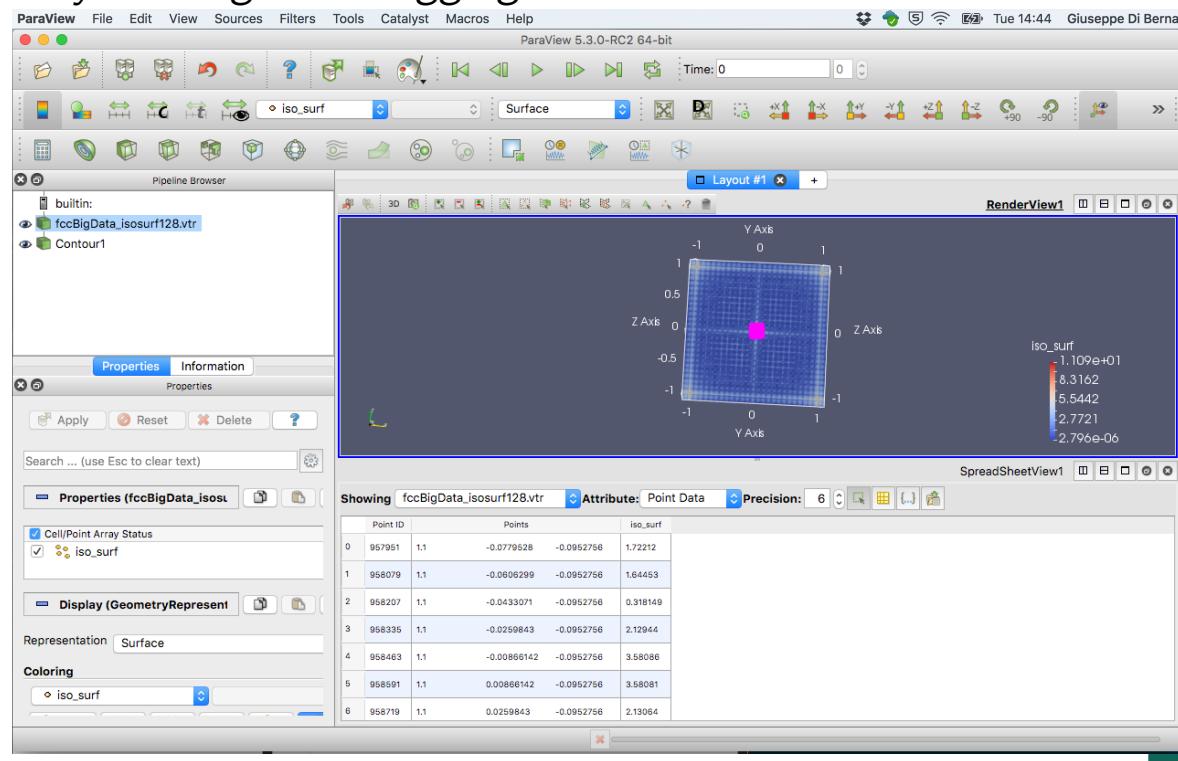
- **Select Blocks:** Select blocks in a multiblock dataset

- **Int. Select Cells**

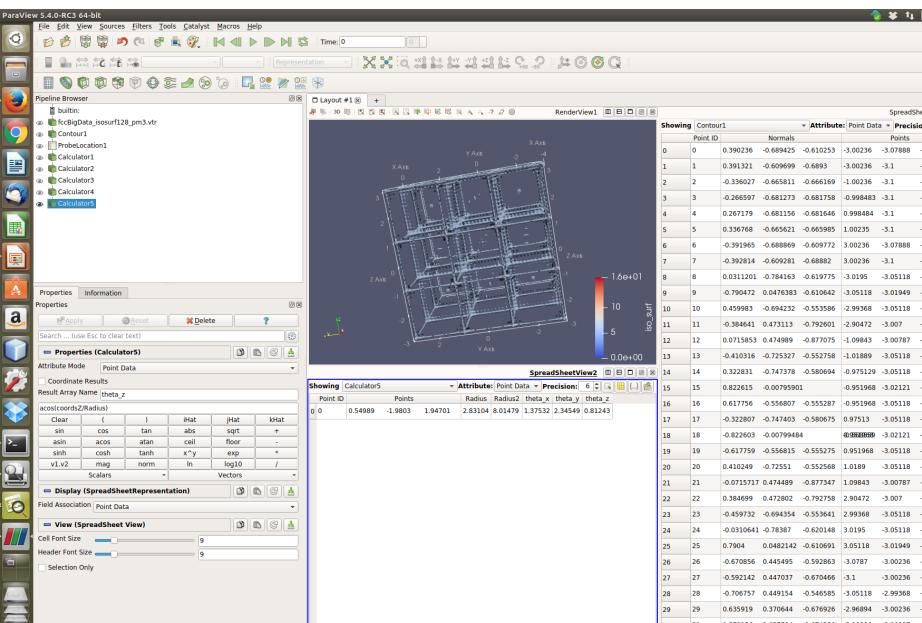
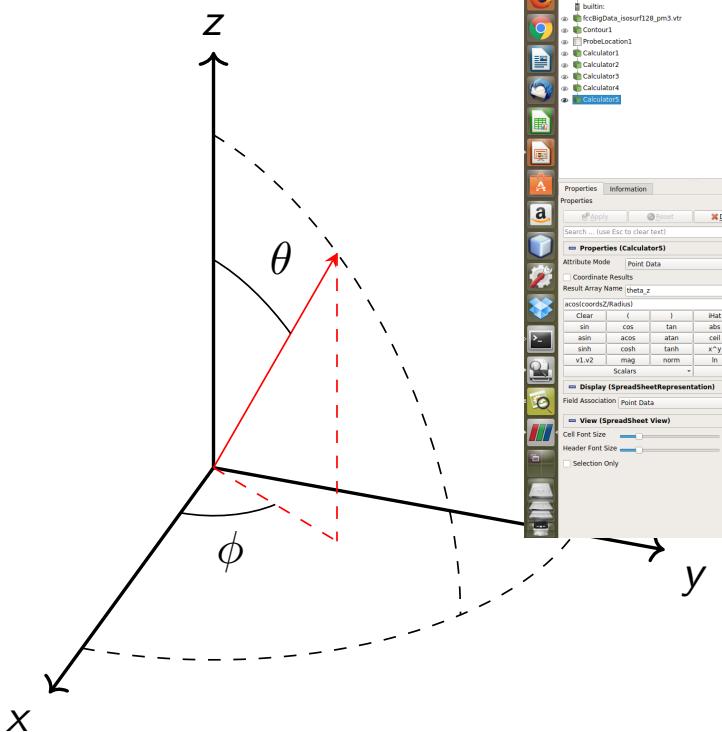
- **Int. Select Points**

Query Data Spatially - Selection

...by clicking and dragging the mouse in the 3D view



Hands-On: extract the tuple of radius and angles



The End

